



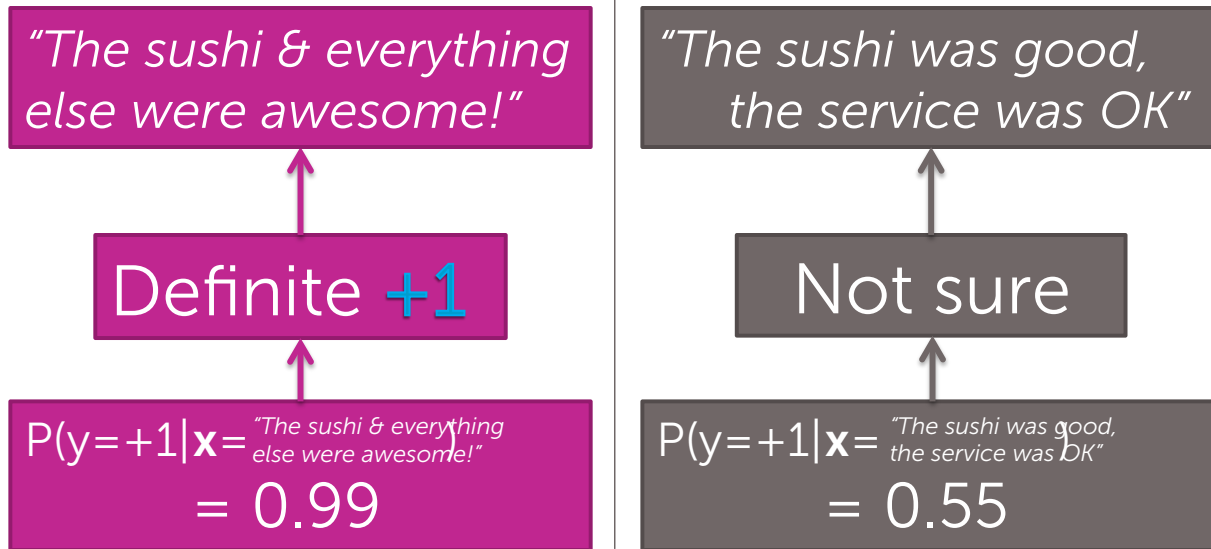
Linear classifiers:



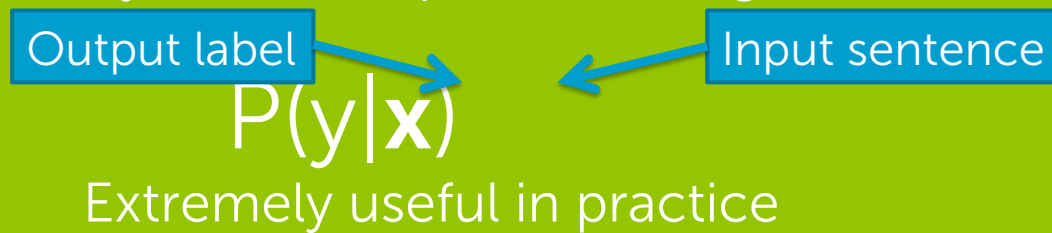
Parameter learning

Emily Fox & Carlos Guestrin  
Machine Learning Specialization  
University of Washington

# Learn a probabilistic classification model



Many classifiers provide a degree of certainty:

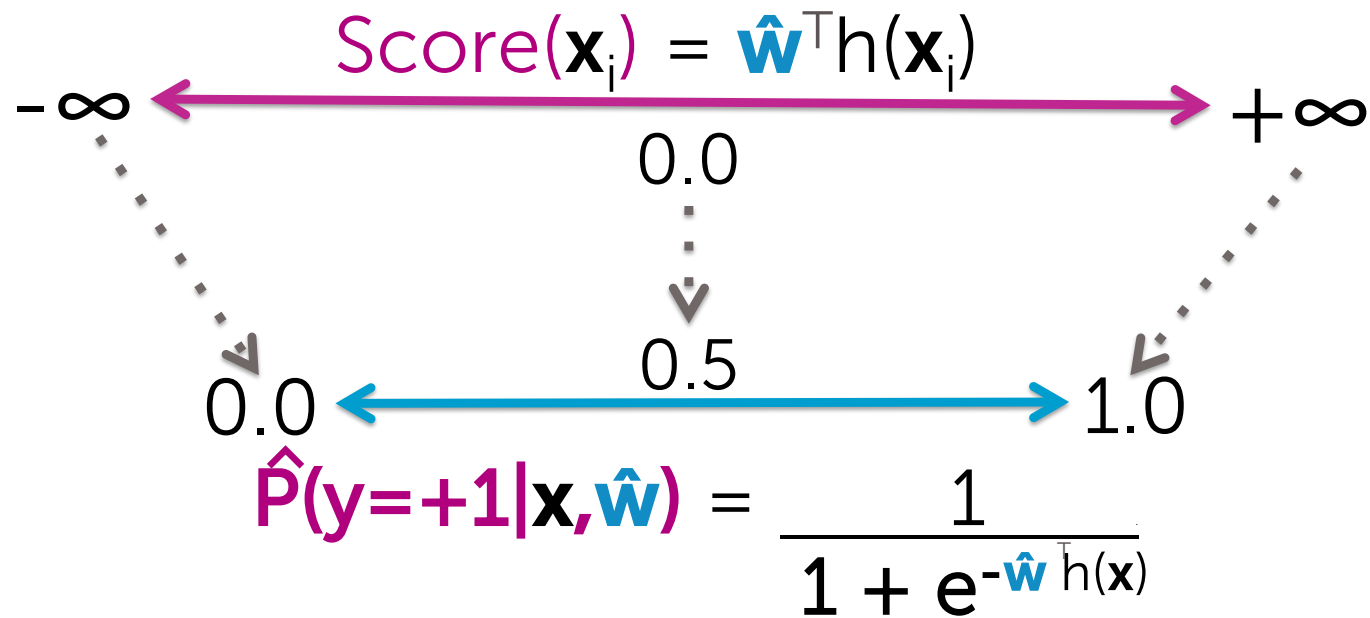


# A (linear) classifier

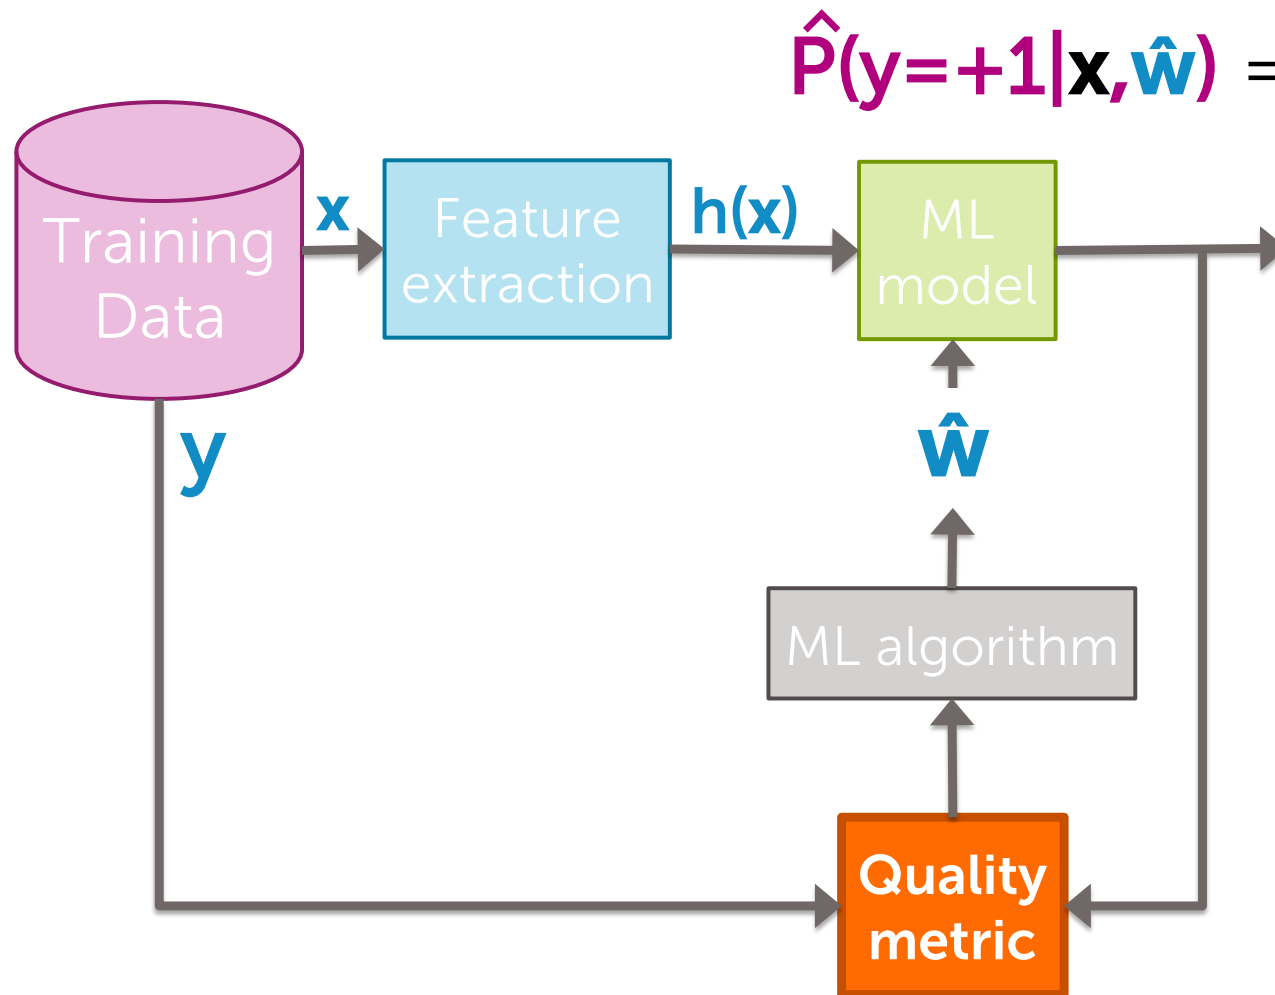
- Will use training data to learn a weight or coefficient for each word

Word	Coefficient	Value
	$\hat{w}_0$	-2.0
good	$\hat{w}_1$	1.0
great	$\hat{w}_2$	1.5
awesome	$\hat{w}_3$	2.7
bad	$\hat{w}_4$	-1.0
terrible	$\hat{w}_5$	-2.1
awful	$\hat{w}_6$	-3.3
restaurant, the, we, ...	$\hat{w}_7, \hat{w}_8, \hat{w}_9, \dots$	0.0
...		...

# Logistic regression model



# Quality metric for logistic regression: Maximum likelihood estimation



$$\hat{P}(y=+1|\mathbf{x},\hat{\mathbf{w}}) = \frac{1}{1 + e^{-\hat{\mathbf{w}}^T h(\mathbf{x})}}$$

# Learning problem

Training data:

$N$  observations  $(\mathbf{x}_i, y_i)$

$x[1] = \text{\#awesome}$	$x[2] = \text{\#awful}$	$y = \text{sentiment}$
2	1	+1
0	2	-1
3	3	-1
4	1	+1
1	1	+1
2	4	-1
0	3	-1
0	1	-1
2	1	+1





# MOVE TO HEAD SHOT



# Finding best coefficients

$x[1] = \text{\#awesome}$	$x[2] = \text{\#awful}$	$y = \text{sentiment}$
2	1	+1
0	2	-1
3	3	-1
4	1	+1
1	1	+1
2	4	-1
0	3	-1
0	1	-1
2	1	+1

# Finding best coefficients

x[1] = #awesome	x[2] = #awful	y = sentiment
0	2	-1
3	3	-1
2	4	-1
0	3	-1
0	1	-1
2	4	-1
0	3	-1
0	1	-1

x[1] = #awesome	x[2] = #awful	y = sentiment
2	1	+1
4	1	+1
1	1	+1
2	1	+1
1	1	+1
2	1	+1

# Finding best coefficients

x[1] = #awesome	x[2] = #awful	y = sentiment
0	2	-1
3	3	-1
2	4	-1
0	3	-1
0	1	-1

$$P(y=+1|\mathbf{x}_i, \mathbf{w}) = 0.0$$

x[1] = #awesome	x[2] = #awful	y = sentiment
2	1	+1
4	1	+1
1	1	+1
2	1	+1

$$P(y=+1|\mathbf{x}_i, \mathbf{w}) = 1.0$$

Pick  $\hat{\mathbf{w}}$  that makes

Quality metric = Likelihood function

Negative data points

$$P(y=+1|\mathbf{x}_i, \mathbf{w}) = 0.0$$

Positive data points

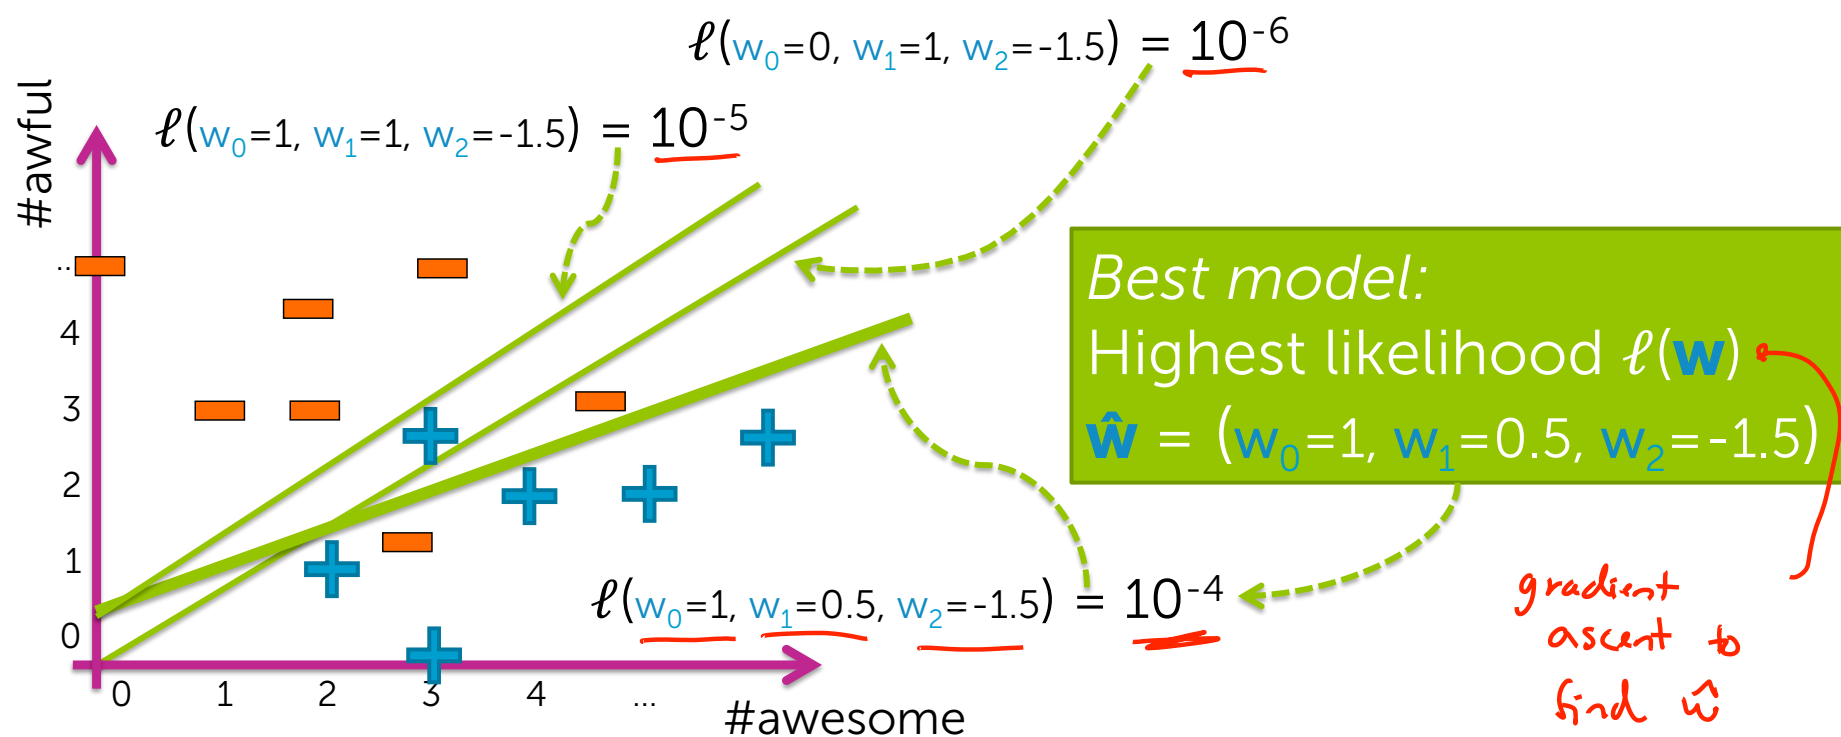
$$P(y=+1|\mathbf{x}_i, \mathbf{w}) = 1.0$$

No  $\hat{\mathbf{w}}$  achieves perfect predictions (usually)

**Likelihood**  $\ell(\mathbf{w})$ : Measures quality of fit for model with coefficients  $\mathbf{w}$

# Find "best" classifier

Maximize likelihood over all possible  $w_0, w_1, w_2$





## Data likelihood

# Quality metric: probability of data

$x[1] = \text{\#awesome}$	$x[2] = \text{\#awful}$	$y = \text{sentiment}$
2	1	+1

$x_1$

$y_1$

If model good, should predict:

$\hat{y}_1 = +1$

Pick  $w$  to maximize:

$$P(y = +1 | x_1, w) = P(y = +1 | x[1]=2, x[2]=1, w)$$

$x[1] = \text{\#awesome}$	$x[2] = \text{\#awful}$	$y = \text{sentiment}$
0	2	-1

$x_2$

$y_2$

If model good, should predict:

$\hat{y}_2 = -1$

Pick  $w$  to maximize:

$$P(y = -1 | x_2, w)$$

# Maximizing likelihood (probability of data)

Data point	x[1]	x[2]	y	Choose $w$ to maximize
$x_1, y_1$	2	1	+1	$P(y=+1 x_1, w) = P(y=+1 x[1]=2, x[2]=1, w)$
$x_2, y_2$	0	2	-1	$P(y=-1 x_2, w)$
$x_3, y_3$	3	3	<u>-1</u>	$P(y=-1 x_3, w)$
$x_4, y_4$	4	1	<u>+1</u>	$P(y=+1 x_4, w)$
$x_5, y_5$	1	1	+1	
$x_6, y_6$	2	4	-1	
$x_7, y_7$	0	3	-1	
$x_8, y_8$	0	1	-1	
$x_9, y_9$	2	1	+1	

Must combine into single measure of quality ?

Multiply probabilities

$$P(y=+1|x_1, w) P(y=-1|x_2, w) P(y=-1|x_3, w) \dots$$

The reason you multiply is that you assume that every row is independent of each other.



# Learn logistic regression model with maximum likelihood estimation (MLE)


Data point	x[1]	x[2]	y	Choose $\mathbf{w}$ to maximize
$\mathbf{x}_1, y_1$	2	1	<u><math>y_1 = +1</math></u>	$P(\underline{y=+1}   \mathbf{x}[1]=2, \mathbf{x}[2]=1, \mathbf{w})$
$\mathbf{x}_2, y_2$	0	2	<u><math>-1</math></u>	$P(\underline{y=-1}   \mathbf{x}[1]=0, \mathbf{x}[2]=2, \mathbf{w})$
$\mathbf{x}_3, y_3$	3	3	$-1$	$P(y=-1   \mathbf{x}[1]=3, \mathbf{x}[2]=3, \mathbf{w})$
$\mathbf{x}_4, y_4$	4	1	$+1$	$P(y=+1   \mathbf{x}[1]=4, \mathbf{x}[2]=1, \mathbf{w})$

$$\ell(\mathbf{w}) = \underbrace{P(y_1 | \mathbf{x}_1, \mathbf{w})}_{\text{P}(y_1|\mathbf{x}_1, \mathbf{w})} \underbrace{P(y_2 | \mathbf{x}_2, \mathbf{w})}_{\text{P}(y_2|\mathbf{x}_2, \mathbf{w})} \underbrace{P(y_3 | \mathbf{x}_3, \mathbf{w})}_{\text{P}(y_3|\mathbf{x}_3, \mathbf{w})} \underbrace{P(y_4 | \mathbf{x}_4, \mathbf{w})}_{\text{P}(y_4|\mathbf{x}_4, \mathbf{w})}$$

Num. of data points  $\rightarrow N$

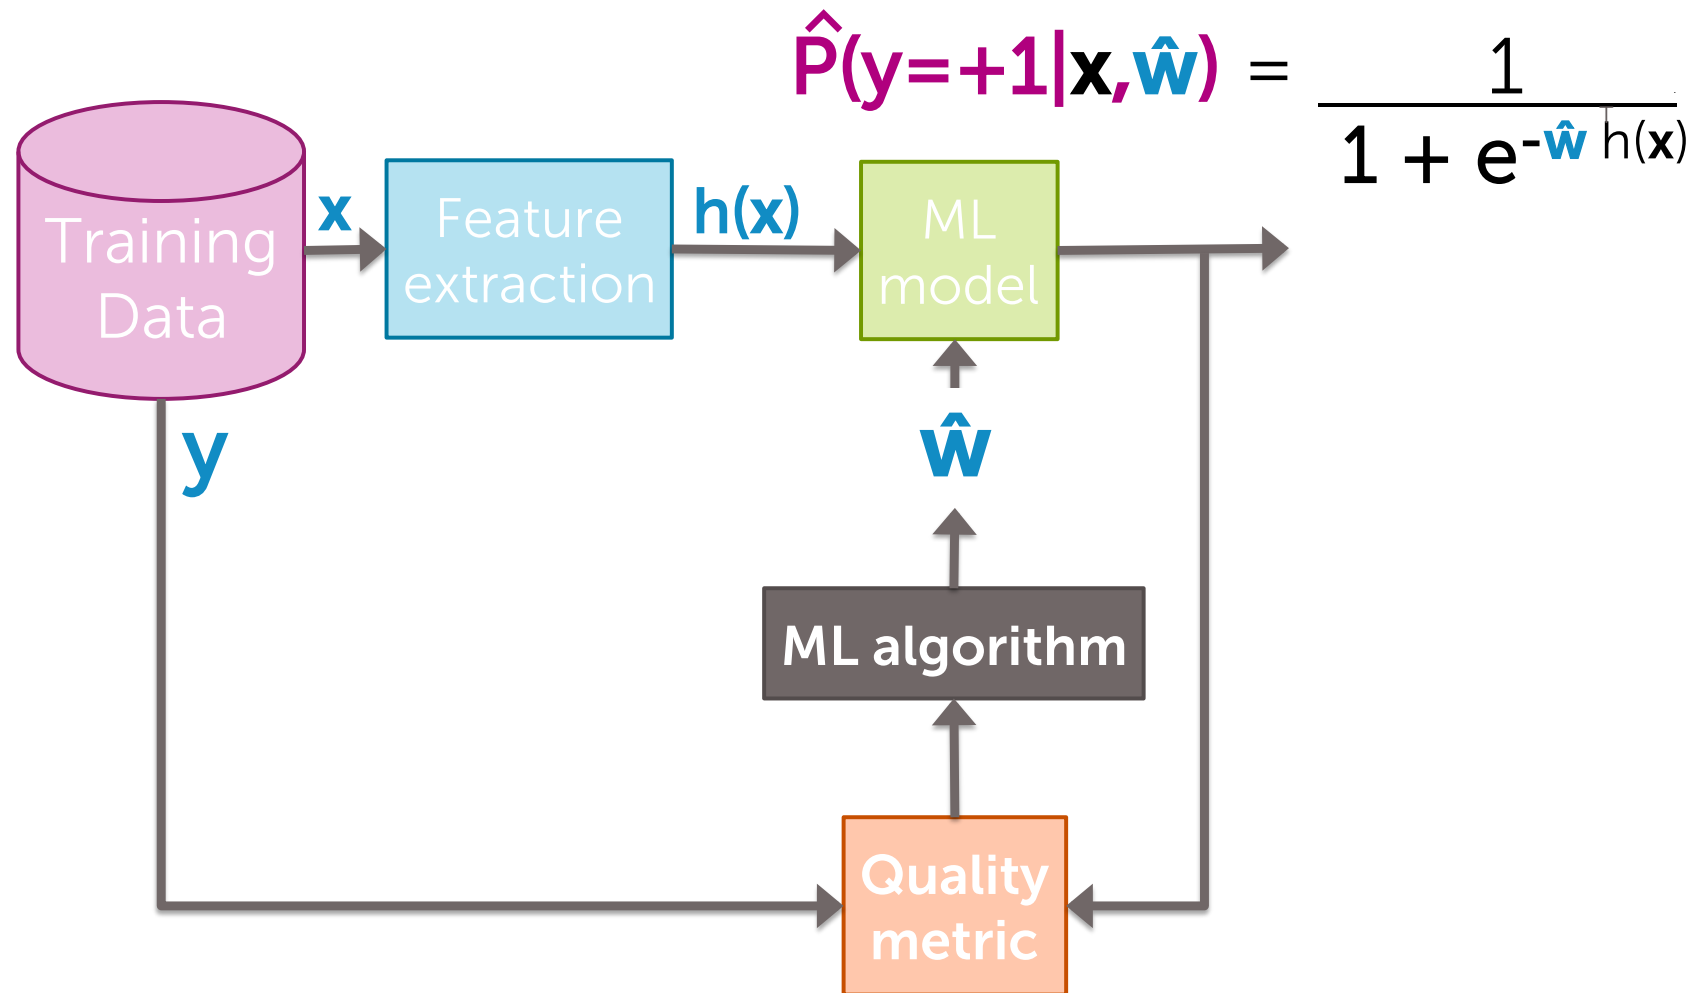
$$\ell(\mathbf{w}) = \prod_{i=1}^N P(y_i | \mathbf{x}_i, \mathbf{w})$$

pick  $\mathbf{w}$  to make this fn. as large as possible



# MOVE TO FULL BODY SHOT

# Finding best linear classifier with gradient ascent





# MOVE TO HEAD SHOT

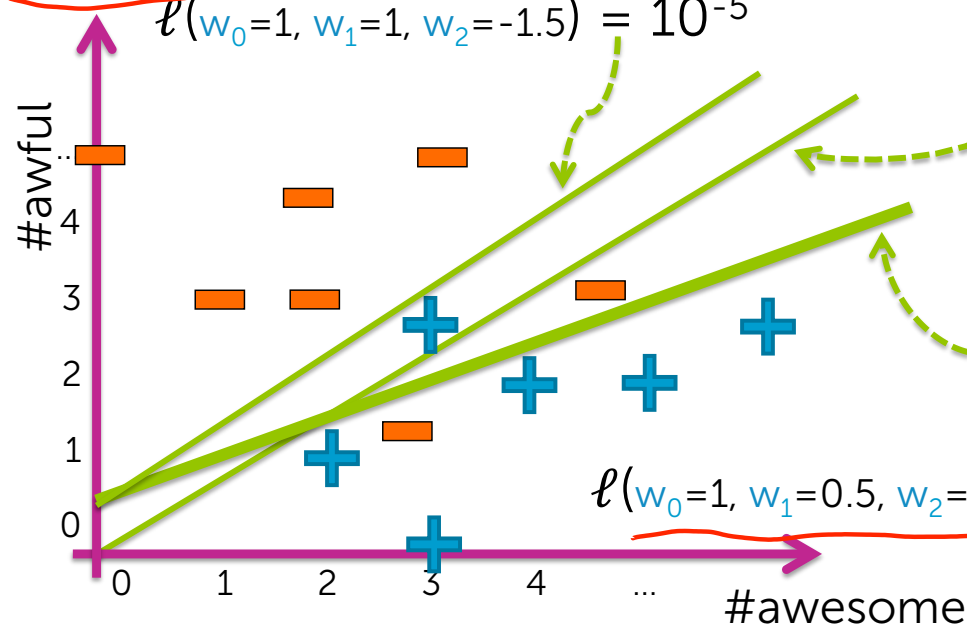
# Find "best" classifier

Maximize likelihood over all possible  $w_0, w_1, w_2$

$$\ell(\mathbf{w}) = \prod_{i=1}^N P(y_i | \mathbf{x}_i, \mathbf{w})$$

$$\ell(w_0=0, w_1=1, w_2=-1.5) = 10^{-6}$$

$$\ell(w_0=1, w_1=1, w_2=-1.5) = 10^{-5}$$



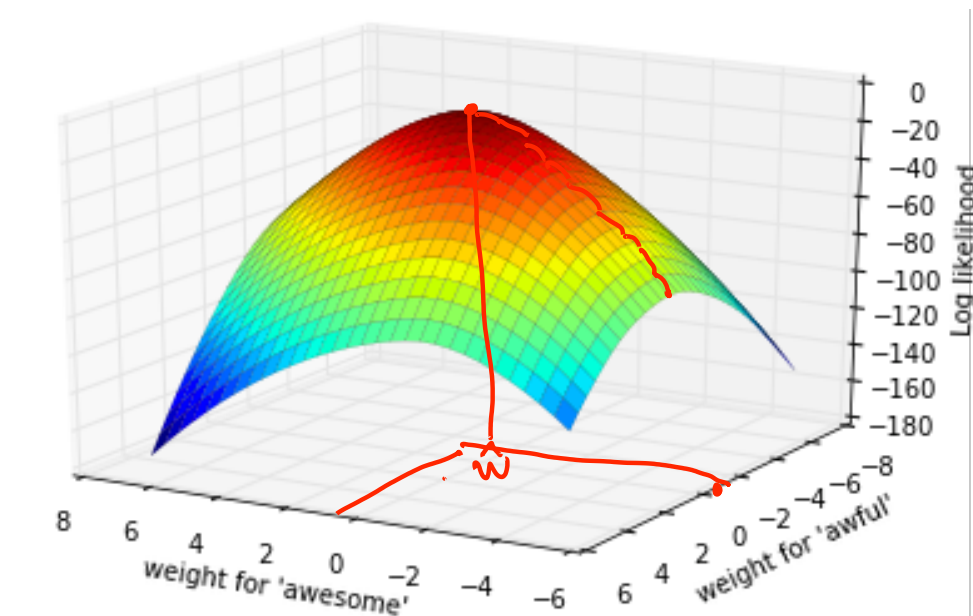
Best model:

Highest likelihood  $\ell(\mathbf{w})$

$$\hat{\mathbf{w}} = (w_0=1, w_1=0.5, w_2=-1.5)$$

optimize with  
gradient ascent

# Maximizing likelihood




No closed-form solution → use gradient ascent

Maximize function over all possible  $w_0, w_1, w_2$

$$\max_{w_0, w_1, w_2} \prod_{i=1}^N P(y_i \mid \mathbf{x}_i, \mathbf{w})$$

$\ell(w_0, w_1, w_2)$  is a function of 3 variables



# MOVE TO FULL BODY SHOT



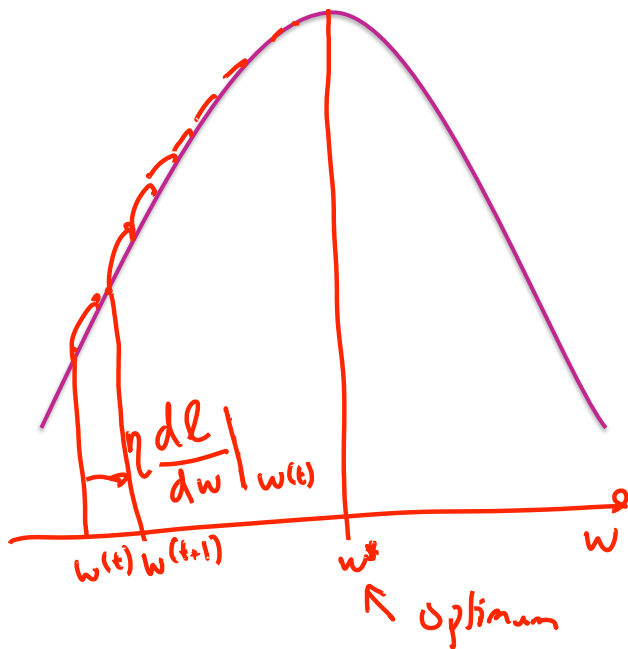


## Review of gradient ascent



# MOVE TO HEAD SHOT

# Finding the max via hill climbing



Algorithm:

**while** not converged

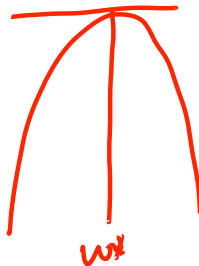
$$w^{(t+1)} \leftarrow w^{(t)} + \eta \frac{d\ell}{dw} \bigg|_{w^{(t)}}$$

step size

# Convergence criteria

For convex functions,  
optimum occurs when

$$\frac{d\ell}{dw} = 0$$



In practice, stop when

$$\left. \frac{d\ell}{dw} \right|_{w^{(t)}} < \epsilon$$

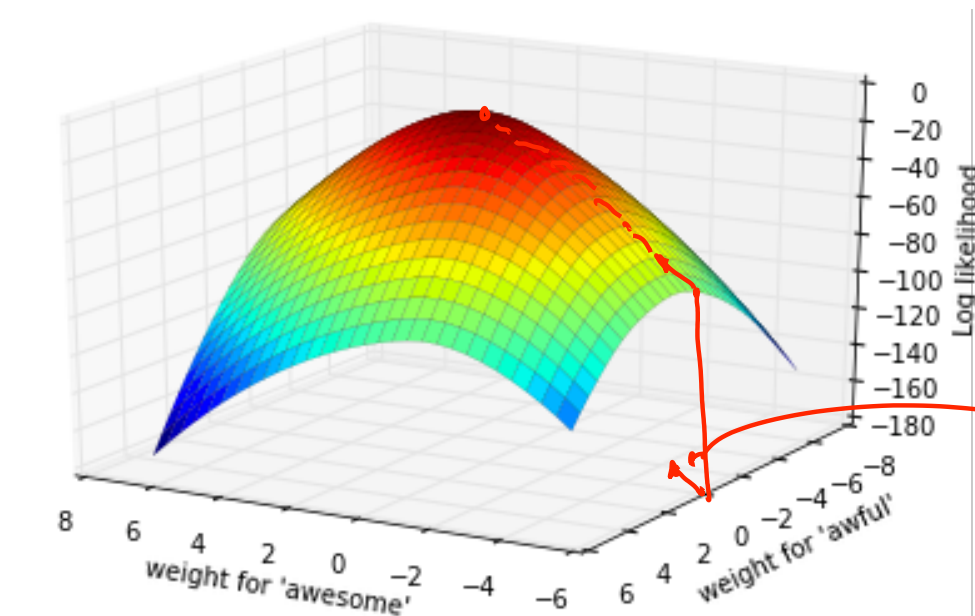
↑  
tolerance

Algorithm:

**while** not converged

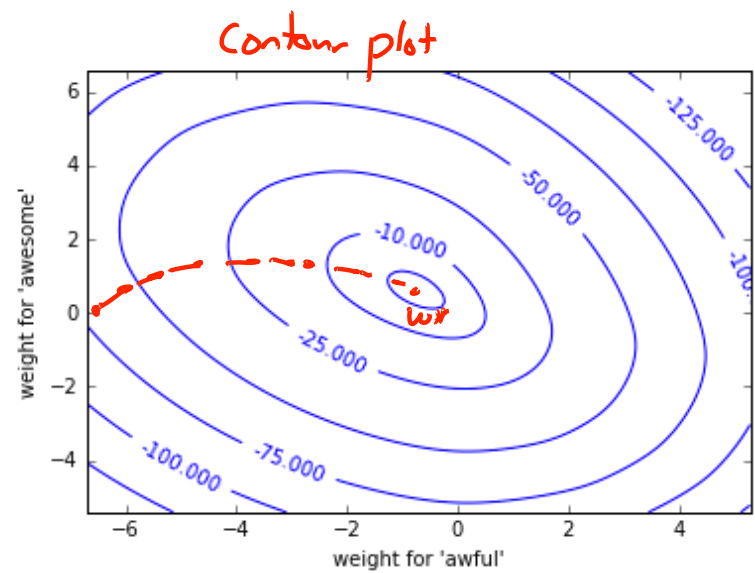
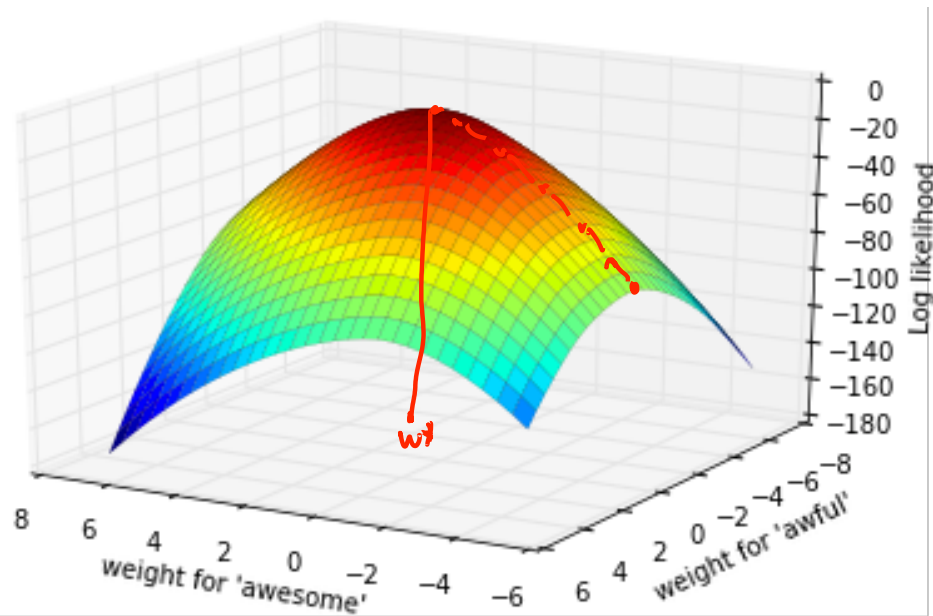
$$w^{(t+1)} \leftarrow w^{(t)} + \eta \left. \frac{d\ell}{dw} \right|_{w^{(t)}}$$

# Moving to multiple dimensions: Gradients

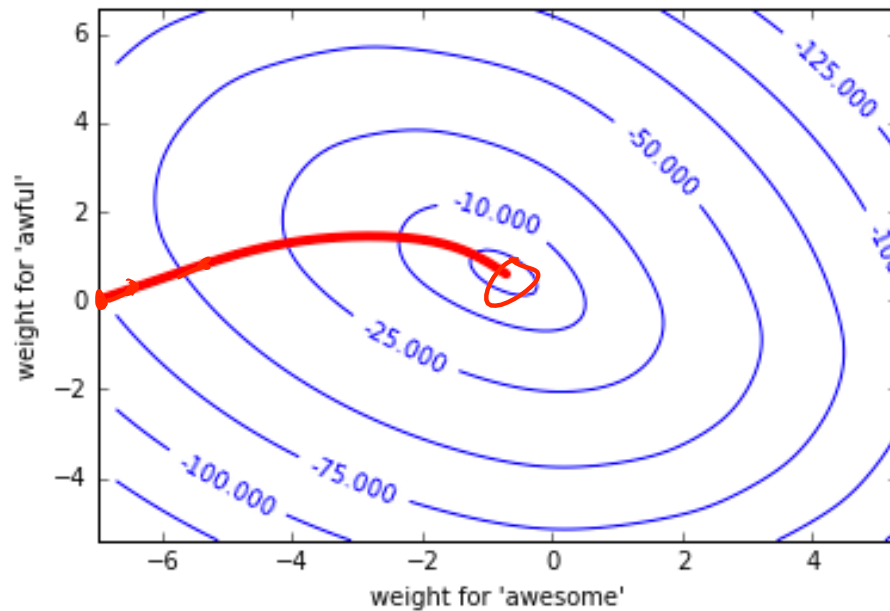


$$\nabla \ell(\mathbf{w}) = \begin{bmatrix} \frac{\partial \ell}{\partial w_0} \\ \frac{\partial \ell}{\partial w_1} \\ \vdots \\ \frac{\partial \ell}{\partial w_D} \end{bmatrix} \leftarrow D+1 \text{ dim vector}$$

# Contour plots



# Gradient ascent




Algorithm:

$w^{(0)} = 0$  , random , or something smart.

**while** not converged

$$\mathbf{w}^{(t+1)} \leftarrow \mathbf{w}^{(t)} + \eta \nabla \ell(\mathbf{w}^{(t)})$$

step size



# MOVE TO FULL BODY SHOT





## Learning algorithm for logistic regression



# MOVE TO HEAD SHOT

# Derivative of (log-)likelihood

Sum over data points

Feature value

Difference between truth and prediction

$$\frac{\partial \ell(\mathbf{w})}{\partial \mathbf{w}_j} = \sum_{i=1}^N h_j(\mathbf{x}_i) \left( \mathbb{1}[y_i = +1] - \underbrace{P(y = +1 \mid \mathbf{x}_i, \mathbf{w})}_{\text{predict } x_i \text{ is positive}} \right)$$

$j$  ranges from  $0 \sim D$ , num of features  
 $i$  ranges from  $1 \sim N$ , num of data points

- output = 1 if  $y_i$  positive  
- output = 0 if  $y_i$  negative  
Indicator function:

$$\mathbb{1}[y_i = +1] = \begin{cases} 1 & \text{if } y_i = +1 \\ 0 & \text{if } y_i = -1 \end{cases}$$

# Computing derivative

$$\frac{\partial \ell(\mathbf{w}^{(t)})}{\partial \mathbf{w}_j} = \sum_{i=1}^N h_j(\mathbf{x}_i) \left( \mathbb{1}[y_i = +1] - P(y = +1 \mid \mathbf{x}_i, \mathbf{w}^{(t)}) \right)$$

$\mathbf{w}^{(t)}$ :

$w_0^{(t)}$	0
$w_1^{(t)}$	1
$w_2^{(t)}$	-2

$$\frac{\partial \ell}{\partial w_1}$$

$h_1(x) = \# \text{ awesome}$  parameter  $w_1$  multiplies 1st feature  $h_1(x)$

x[1]	x[2]	y	P(y=+1 x,w)	Contribution to derivative for $w_1$
2	1	+1	0.5	$2(1 - 0.5) = 1$
0	2	-1	0.02	$0(0 - 0.02) = 0$
3	3	-1	0.05	$3(0 - 0.05) = -0.15$
4	1	+1	0.88	$4(1 - 0.88) = 0.48$

Total derivative:

$$\frac{\partial \ell(\mathbf{w}^{(t)})}{\partial w_1} = 1 + 0 - 0.15 + 0.48 = 1.33$$

$$w_1^{(t+1)} = w_1^{(t)} + \eta \frac{\partial \ell(\mathbf{w}^{(t)})}{\partial w_1} \quad | \quad \eta = 0.1$$

$$= 1 + 0.1 \times 1.33 = 1.133$$

# Derivative of (log-)likelihood: Interpretation

Sum over data points

Feature value

Difference between truth and prediction

$$\frac{\partial \ell(\mathbf{w})}{\partial \mathbf{w}_j} = \sum_{i=1}^N h_j(\mathbf{x}_i) \left( \mathbb{1}[y_i = +1] - P(y = +1 | \mathbf{x}_i, \mathbf{w}) \right)$$

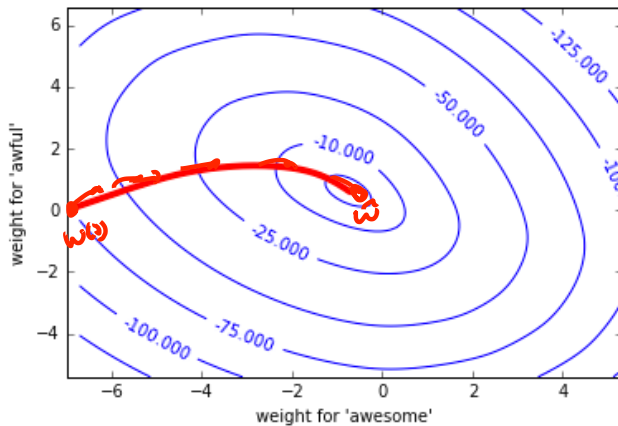
$\Delta_i$

If  $h_j(\mathbf{x}_i) = 1$ :

	$P(y=+1 \mathbf{x}_i, \mathbf{w}) \approx 1$	$P(y=+1 \mathbf{x}_i, \mathbf{w}) \approx 0$
$y_i = +1$	$\Delta_i = (1 - 1) \approx 0$ $\hookrightarrow$ don't change anything!	$\Delta_i \approx 1 \Rightarrow$ increase $w_j$ $\Rightarrow$ Score( $\mathbf{x}_i$ ) becomes larger $\Rightarrow P(y=+1 \mathbf{x}_i, \mathbf{w})$ increases
$y_i = -1$	$\Delta_i = -1 \Rightarrow w_j$ to decrease $\Rightarrow$ Score( $\mathbf{x}_i$ ) decreases $\Rightarrow P(y=+1 \mathbf{x}_i, \mathbf{w})$ decrease	$\Delta_i \approx 0$ $\Rightarrow$ don't change anything

increase parameter  $w$   
 $\Rightarrow$  increase score  
 $\Rightarrow$  increase probability

# Summary of gradient ascent for logistic regression



init  $\mathbf{w}^{(1)} = 0$  (or randomly, or smartly),  $t = 1$

while  $\|\nabla \ell(\mathbf{w}^{(t)})\| > \epsilon$


for  $j = 0, \dots, D$

$$\text{partial}[j] = \sum_{i=1}^N h_j(\mathbf{x}_i) \left( \mathbb{1}[y_i = +1] - \underbrace{P(y = +1 \mid \mathbf{x}_i, \mathbf{w}^{(t)})}_{\frac{1}{1 + e^{-\mathbf{w}^{(t)} \cdot \mathbf{h}(\mathbf{x}_i)}}} \right)$$

$$\mathbf{w}_j^{(t+1)} \leftarrow \mathbf{w}_j^{(t)} + \eta \text{partial}[j]$$

$$t \leftarrow t + 1$$

step size  $\uparrow$   $\frac{\partial \ell(\mathbf{w}^{(t)})}{\partial \mathbf{w}_j}$



# MOVE TO FULL BODY SHOT



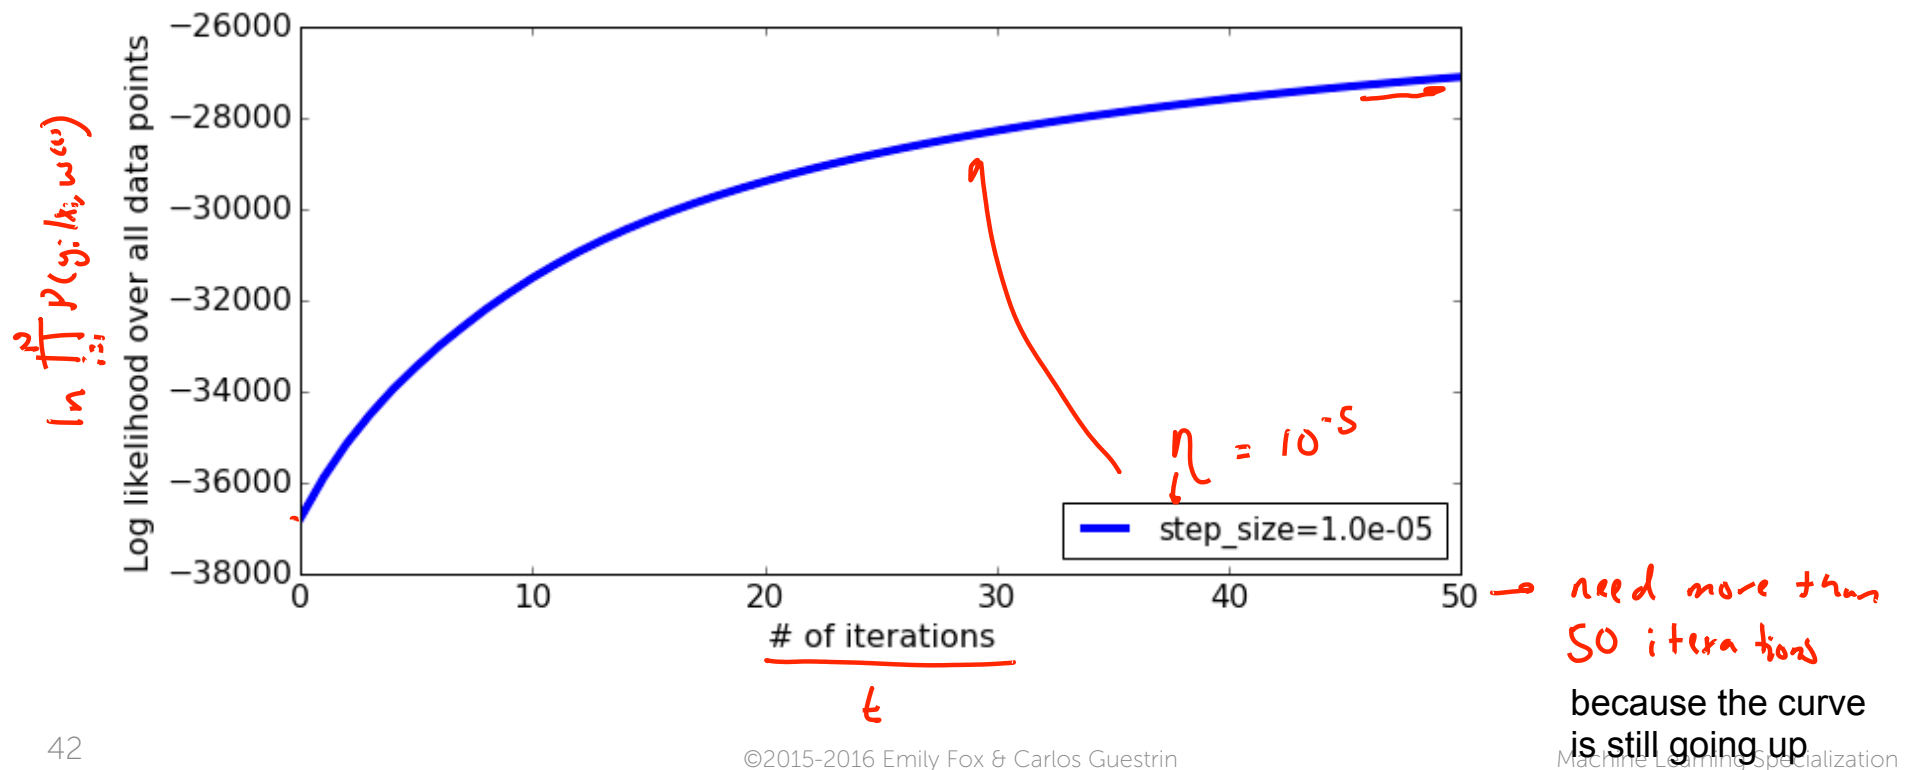
## Choosing the step size $\eta$



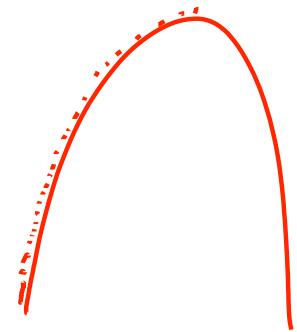
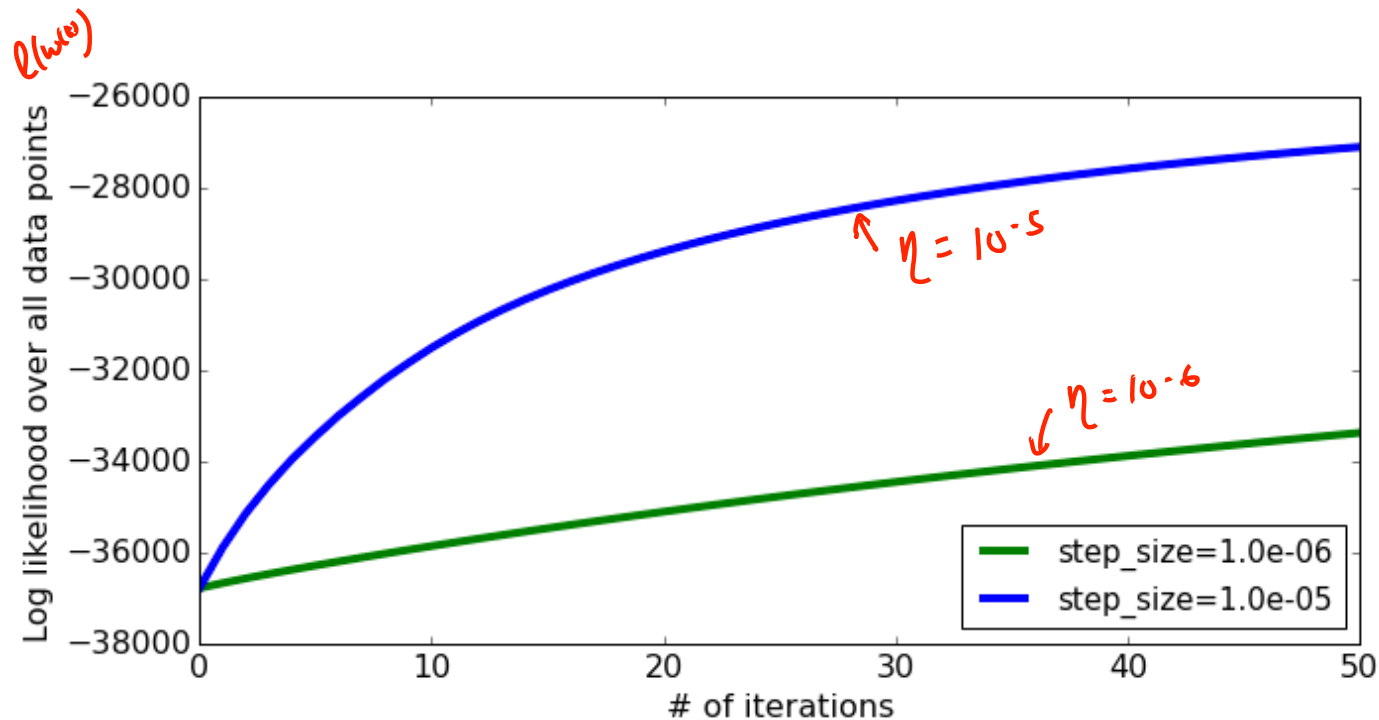


# MOVE TO HEAD SHOT

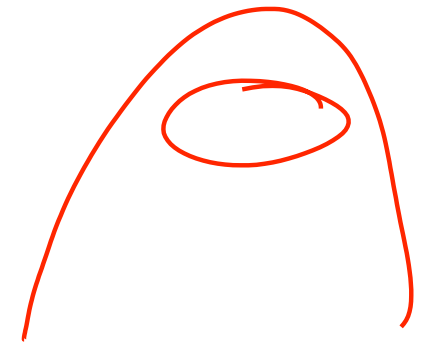
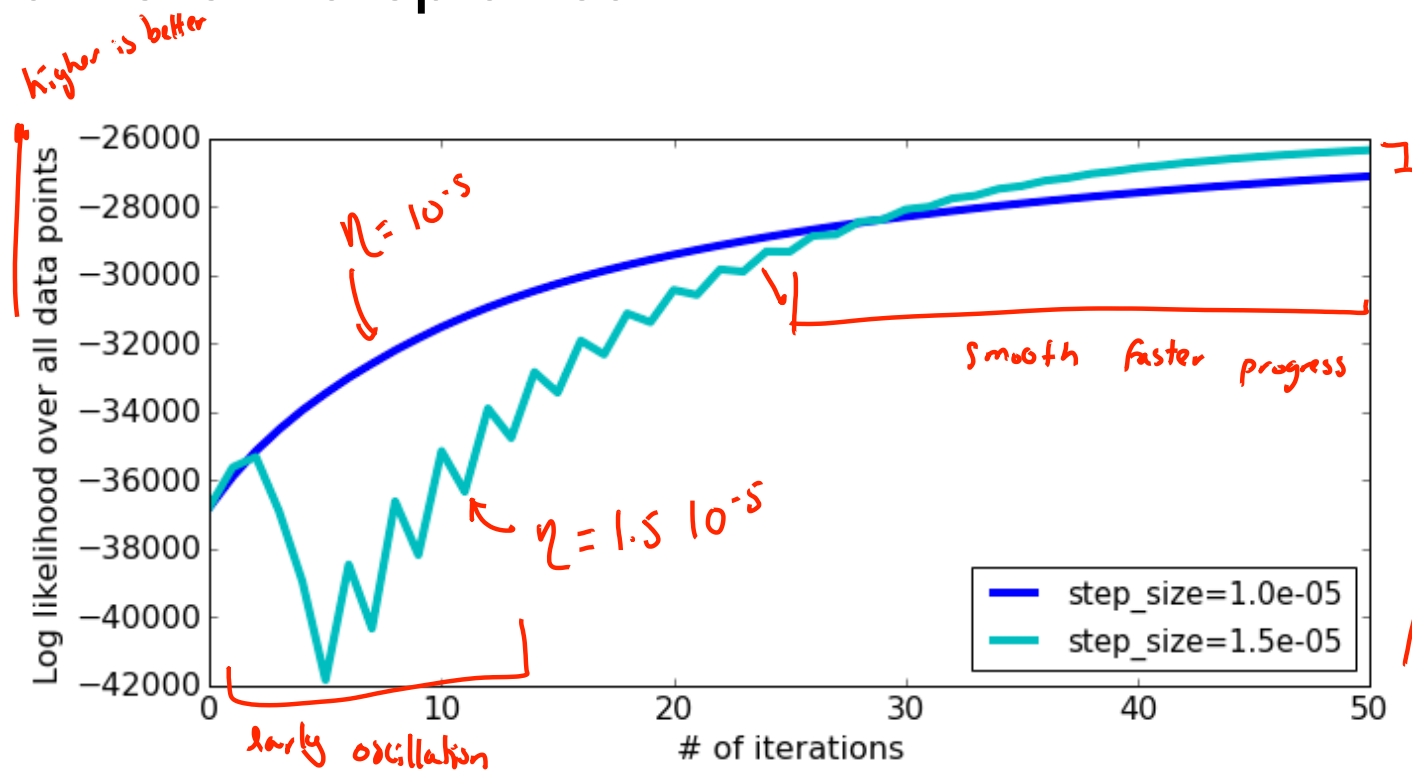
## Learning curve: Plot quality (likelihood) over iterations



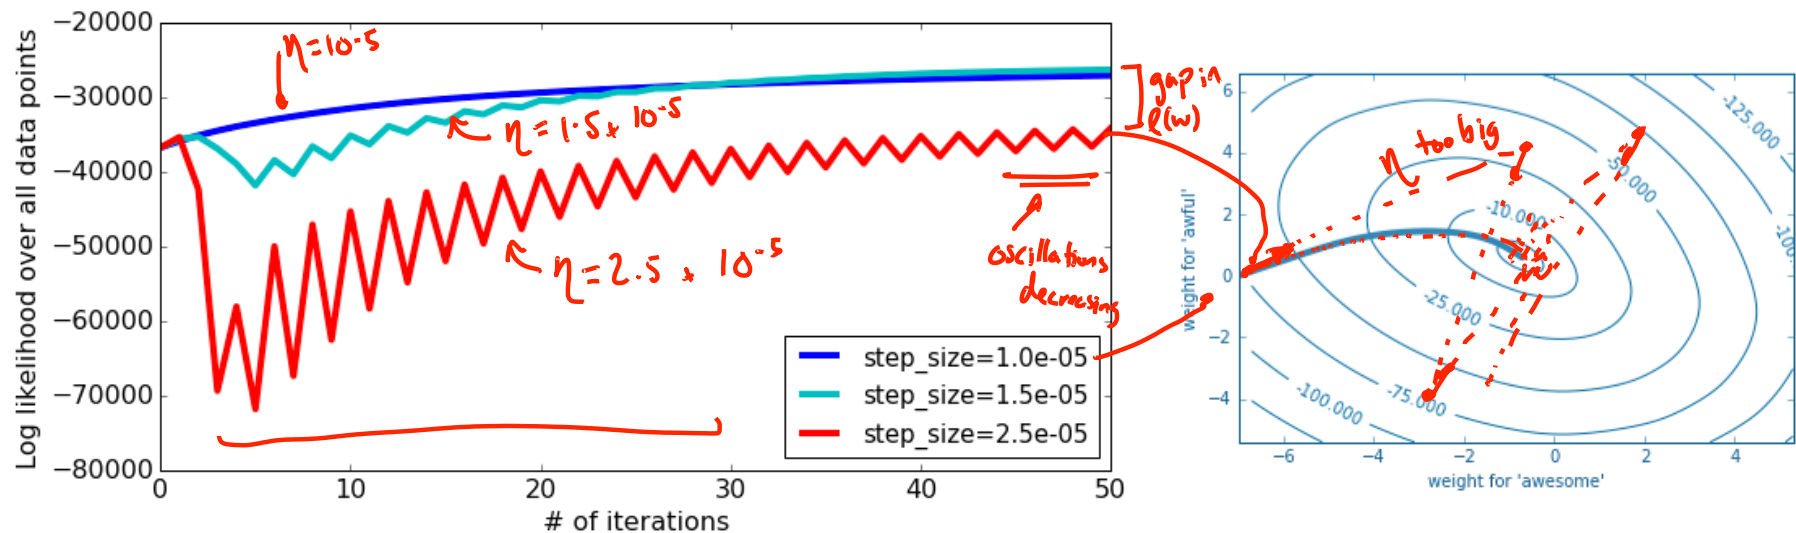
If step size is too small, can take a long time to converge



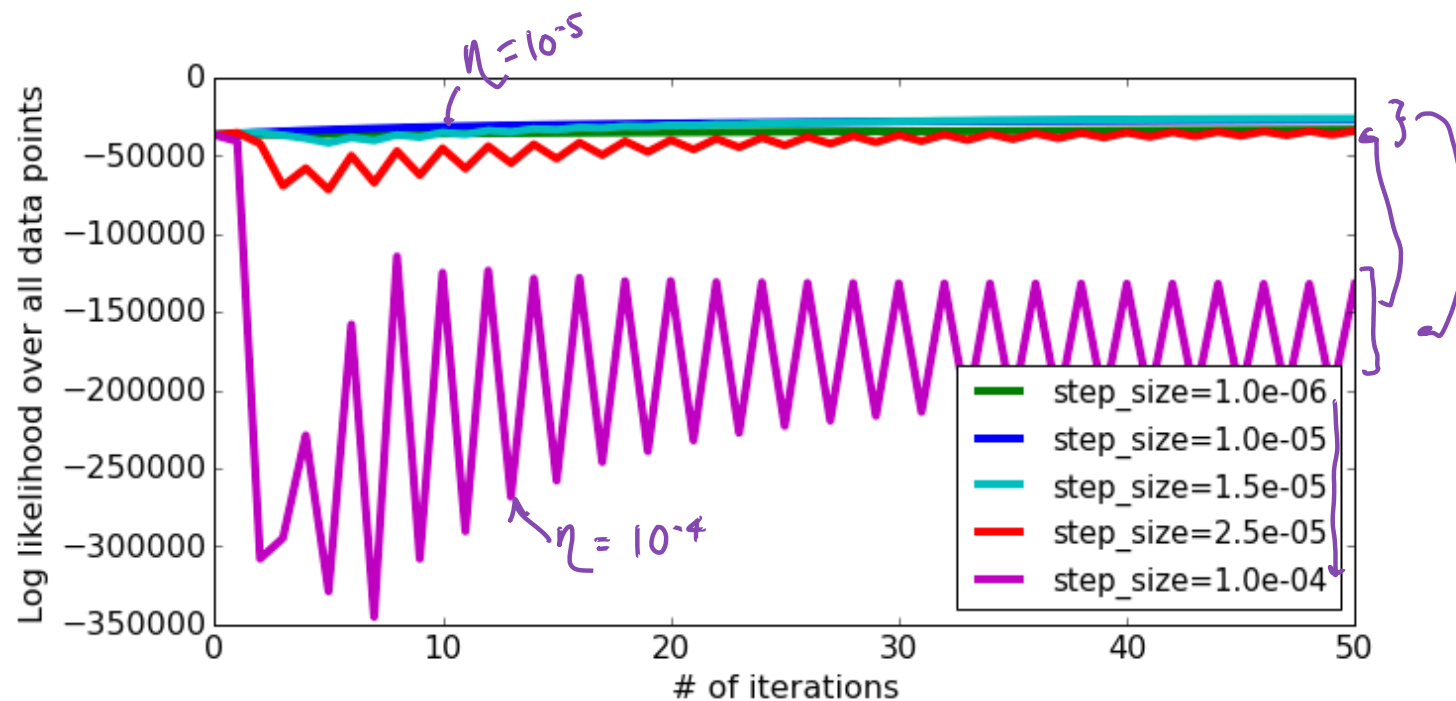
## Compare converge with different step sizes



# Careful with step sizes that are too large



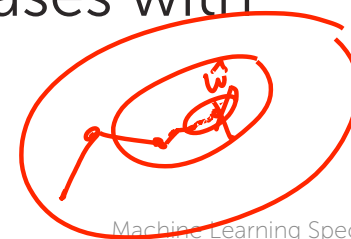
Very large step sizes can even cause divergence or wild oscillations




## Simple rule of thumb for picking step size $\eta$

- Unfortunately, picking step size requires a lot of trial and error ☹
- Try a several values, exponentially spaced
  - **Goal:** plot learning curves to
    - find one  $\eta$  that is too small (smooth but moving too slowly)
    - find one  $\eta$  that is too large (oscillation or divergence)
- Try values in between to find “best”  $\eta$ 
  - ↳ exponentially space, pick one that leads best training data likelihood
- Advanced tip: can also try step size that decreases with iterations, e.g.,

$$\eta_t = \frac{\eta_0}{t}$$





# MOVE TO FULL BODY SHOT



# Deriving the gradient for logistic regression

**VERY  
OPTIONAL**



# MOVE TO HEAD SHOT

## Log-likelihood function

- Goal: choose coefficients  $\mathbf{w}$  maximizing likelihood:

$$\ell(\mathbf{w}) = \prod_{i=1}^N P(y_i \mid \mathbf{x}_i, \mathbf{w})$$

- Math simplified by using log-likelihood – taking (natural) log:

$$\ell\ell(\mathbf{w}) = \ln \prod_{i=1}^N P(y_i \mid \mathbf{x}_i, \mathbf{w})$$

*natural log*

# The log trick, often used in ML...

- Products become sums:

$$\ln a \cdot b = \ln a + \ln b \quad \left| \quad \ln \frac{a}{b} = \ln a - \ln b$$

- Doesn't change maximum!

– If  $\hat{\mathbf{w}}$  maximizes  $f(\mathbf{w})$ :

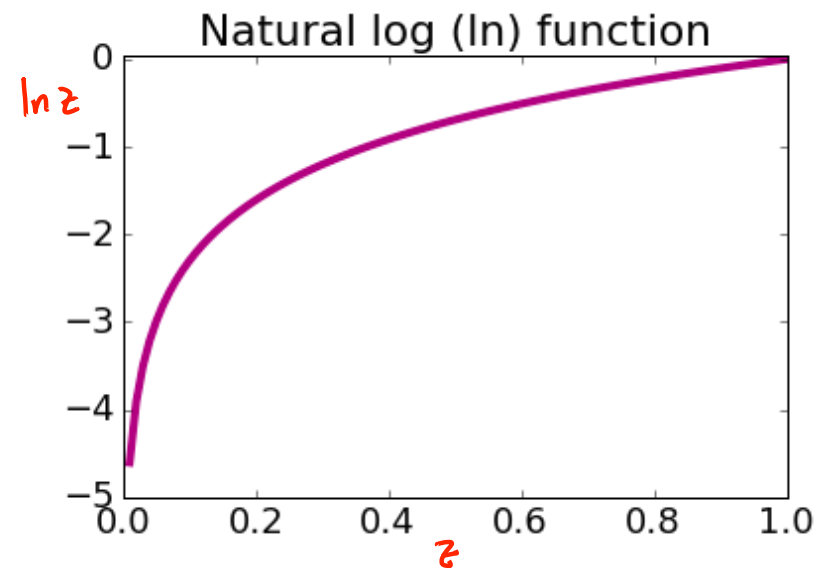
$$\hat{\mathbf{w}} = \arg \max_{\mathbf{w}} f(\mathbf{w})$$


*the  $\mathbf{w}$  that makes  $f(\mathbf{w})$  largest*

– Then  $\hat{\mathbf{w}}_{\ln}$  maximizes  $\ln(f(\mathbf{w}))$ :

$$\hat{\mathbf{w}}_{\ln} = \arg \max_{\mathbf{w}} \ln(f(\mathbf{w}))$$

$$\hat{\mathbf{w}} = \hat{\mathbf{w}}_{\ln}$$





Insert next title slide before Slide 52,  
around 4:55 in  
PL7\_DerivingtheGradient\_1stEdit

## Expressing the log-likelihood

**VERY  
OPTIONAL**

## Using log to turn products into sums

$$\ln \prod_{i=1}^N f_i = \sum_{i=1}^N \ln f_i$$

- The log of the product of likelihoods becomes the sum of the logs:

$$\begin{aligned} \ell\ell(\mathbf{w}) &= \ln \prod_{i=1}^N P(y_i \mid \mathbf{x}_i, \mathbf{w}) \\ &= \sum_{i=1}^N \ln P(y_i \mid \mathbf{x}_i, \mathbf{w}) \end{aligned}$$

# Rewriting log-likelihood

- For simpler math, we'll rewrite likelihood with indicators:

$$\begin{aligned}\ell\ell(\mathbf{w}) &= \sum_{i=1}^N \ln P(y_i \mid \mathbf{x}_i, \mathbf{w}) \\ &= \sum_{i=1}^N [\mathbb{1}[y_i = +1] \ln P(y = +1 \mid \mathbf{x}_i, \mathbf{w}) + \mathbb{1}[y_i = -1] \ln P(y = -1 \mid \mathbf{x}_i, \mathbf{w})]\end{aligned}$$

if  $y_i = +1$

if  $y_i = -1$


✓

○

○

✓





Insert next title slide before Slide 54,  
around 7:33 in  
PL7\_DerivingtheGradient\_1stEdit



Deriving probability that  $y = -1$  given  $x$

**VERY  
OPTIONAL**


## Logistic regression model: $P(y=-1|\mathbf{x}, \mathbf{w})$

- Probability model predicts  $y=+1$ :

$$P(y=+1|\mathbf{x}, \mathbf{w}) = \frac{1}{1 + e^{-\mathbf{w}^T \mathbf{h}(\mathbf{x})}}$$

- Probability model predicts  $y=-1$ :

$$\begin{aligned} P(y=-1|\mathbf{x}, \mathbf{w}) &= 1 - P(y=+1|\mathbf{x}, \mathbf{w}) = 1 - \frac{1}{1 + e^{-\mathbf{w}^T \mathbf{h}(\mathbf{x})}} \\ &= \frac{1 + e^{-\mathbf{w}^T \mathbf{h}(\mathbf{x})} - 1}{1 + e^{-\mathbf{w}^T \mathbf{h}(\mathbf{x})}} = \frac{e^{-\mathbf{w}^T \mathbf{h}(\mathbf{x})}}{1 + e^{-\mathbf{w}^T \mathbf{h}(\mathbf{x})}} \end{aligned}$$



Insert next title slide before Slide 55,  
around 9:15 in  
PL7\_DerivingtheGradient\_1stEdit

## Rewriting the log-likelihood

**VERY  
OPTIONAL**

## Plugging in logistic function for 1 data point

$$P(y = +1 \mid \mathbf{x}, \mathbf{w}) = \frac{1}{1 + e^{-\mathbf{w}^\top h(\mathbf{x})}} \quad P(y = -1 \mid \mathbf{x}, \mathbf{w}) = \frac{e^{-\mathbf{w}^\top h(\mathbf{x})}}{1 + e^{-\mathbf{w}^\top h(\mathbf{x})}}$$

$$\ell\ell(\mathbf{w}) = \mathbb{1}[y_i = +1] \ln P(y = +1 \mid \mathbf{x}_i, \mathbf{w}) + \mathbb{1}[y_i = -1] \ln P(y = -1 \mid \mathbf{x}_i, \mathbf{w})$$

$$= \mathbb{1}[y_i = +1] \ln \frac{1}{1 + e^{-\mathbf{w}^\top h(\mathbf{x}_i)}} + (1 - \mathbb{1}[y_i = +1]) \ln \frac{e^{-\mathbf{w}^\top h(\mathbf{x}_i)}}{1 + e^{-\mathbf{w}^\top h(\mathbf{x}_i)}}$$

$$= -\mathbb{1}[y_i = +1] \ln(1 + e^{-\mathbf{w}^\top h(\mathbf{x}_i)}) + (1 - \mathbb{1}[y_i = +1]) [-\mathbf{w}^\top h(\mathbf{x}_i) - \ln(1 + e^{-\mathbf{w}^\top h(\mathbf{x}_i)})]$$

$$= - (1 - \mathbb{1}[y_i = +1]) \mathbf{w}^\top h(\mathbf{x}_i) - \ln(1 + e^{-\mathbf{w}^\top h(\mathbf{x}_i)})$$

Simpler form


$$\ln e^a = a$$

$$\mathbb{1}[y_i = -1] = 1 - \mathbb{1}[y_i = +1]$$

$$\ln \frac{1}{1 + e^{-\mathbf{w}^\top h(\mathbf{x}_i)}} = -\ln(1 + e^{-\mathbf{w}^\top h(\mathbf{x}_i)})$$

$$\ln \frac{e^{-\mathbf{w}^\top h(\mathbf{x}_i)}}{1 + e^{-\mathbf{w}^\top h(\mathbf{x}_i)}} =$$

$$\underbrace{\ln e^{-\mathbf{w}^\top h(\mathbf{x}_i)}}_{-\mathbf{w}^\top h(\mathbf{x}_i)} - \ln(1 + e^{-\mathbf{w}^\top h(\mathbf{x}_i)})$$



Insert next title slide before Slide 56,  
around 16:56 in  
PL7\_DerivingtheGradient\_1stEdit



## Deriving gradient of log-likelihood

**VERY  
OPTIONAL**



## Gradient for 1 data point

$$\ell(\mathbf{w}) = -(1 - \mathbb{1}[y_i = +1])\mathbf{w}^\top h(\mathbf{x}_i) - \ln(1 + e^{-\mathbf{w}^\top h(\mathbf{x}_i)})$$

$$\frac{\partial \ell}{\partial w_j} = -(1 - \mathbb{1}[y_i = +1]) \frac{\partial w^\top h(\mathbf{x}_i)}{\partial w_j} - \frac{\partial \ln(1 + e^{-w^\top h(\mathbf{x}_i)})}{\partial w_j}$$

$$= -(1 - \mathbb{1}[y_i = +1]) h_j(\mathbf{x}_i) + h_j(\mathbf{x}_i) P(y = -1 | \mathbf{x}_i, \mathbf{w})$$

$$= h_j(\mathbf{x}_i) [\mathbb{1}[y_i = +1] - P(y = +1 | \mathbf{x}_i, \mathbf{w})]$$

$$\begin{aligned} \frac{\partial w^\top h(\mathbf{x}_i)}{\partial w_j} &= h_j(\mathbf{x}_i) \\ \hline \frac{\partial \ln(1 + e^{-w^\top h(\mathbf{x}_i)})}{\partial w_j} &= -h_j(\mathbf{x}_i) \frac{e^{-w^\top h(\mathbf{x}_i)}}{1 + e^{-w^\top h(\mathbf{x}_i)}} \\ &= -h_j(\mathbf{x}_i) P(y = -1 | \mathbf{x}_i, \mathbf{w}) \end{aligned}$$


## Finally, gradient for all data points

- Gradient for one data point:

$$h_j(\mathbf{x}_i) \left( \mathbb{1}[y_i = +1] - P(y = +1 \mid \mathbf{x}_i, \mathbf{w}) \right)$$

- Adding over data points:

$$\frac{\partial \ell}{\partial w_j} = \sum_{i=1}^N h_j(x_i) \left( \mathbb{1}[y_i = +1] - P(y = +1 \mid x_i, w) \right)$$

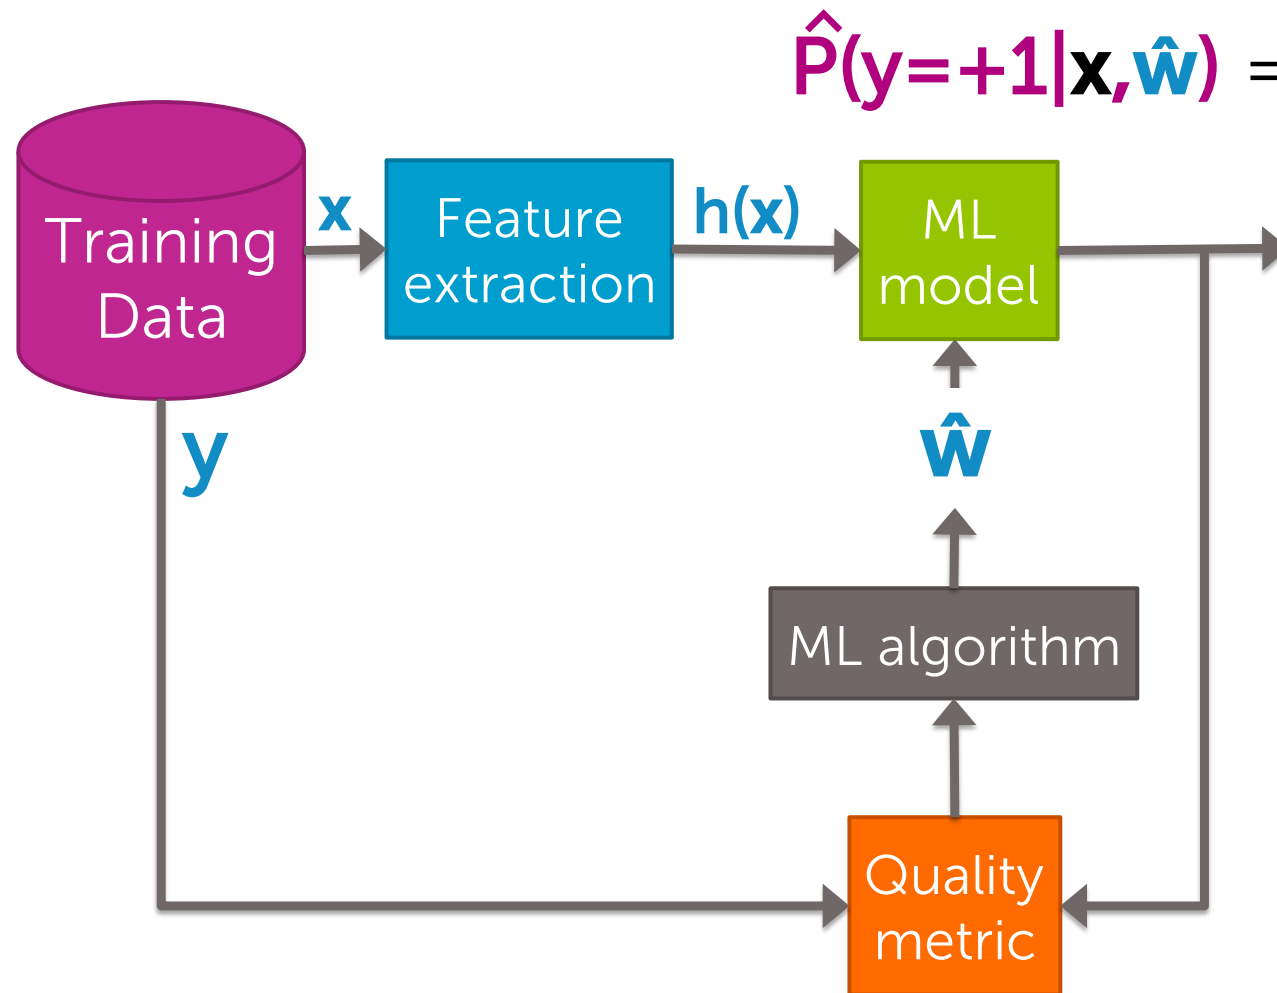


# MOVE TO FULL BODY SHOT

# Summary of logistic regression classifier



# MOVE TO HEAD SHOT



$$\hat{P}(y=+1|\mathbf{x},\hat{\mathbf{w}}) = \frac{1}{1 + e^{-\hat{\mathbf{w}}^T h(\mathbf{x})}}$$



# MOVE TO FULL BODY SHOT

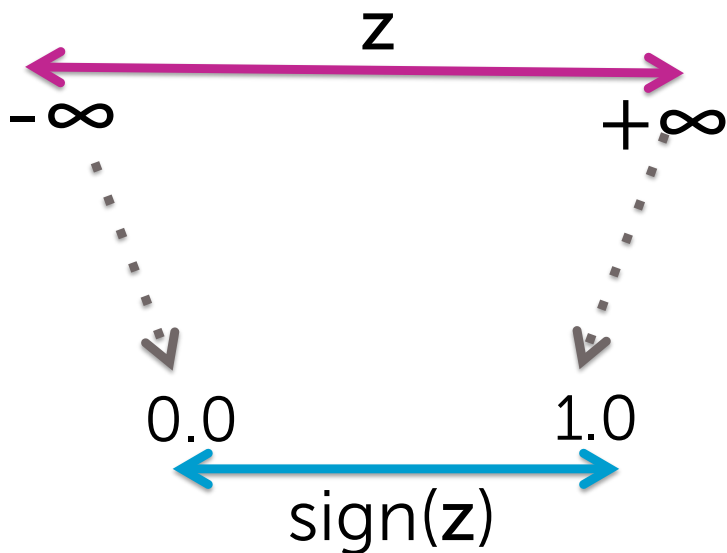


## What you can do now...

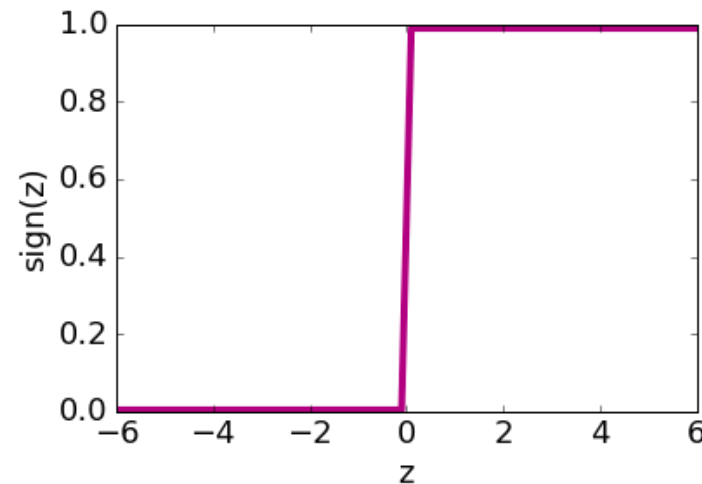
- Measure quality of a classifier using the likelihood function
- Interpret the likelihood function as the probability of the observed data
- Learn a logistic regression model with gradient descent
- (Optional) Derive the gradient descent update rule for logistic regression



# Simplest link function: $\text{sign}(z)$



$$\text{sign}(z) = \begin{cases} +1 & \text{if } z \geq 0 \\ -1 & \text{otherwise} \end{cases}$$

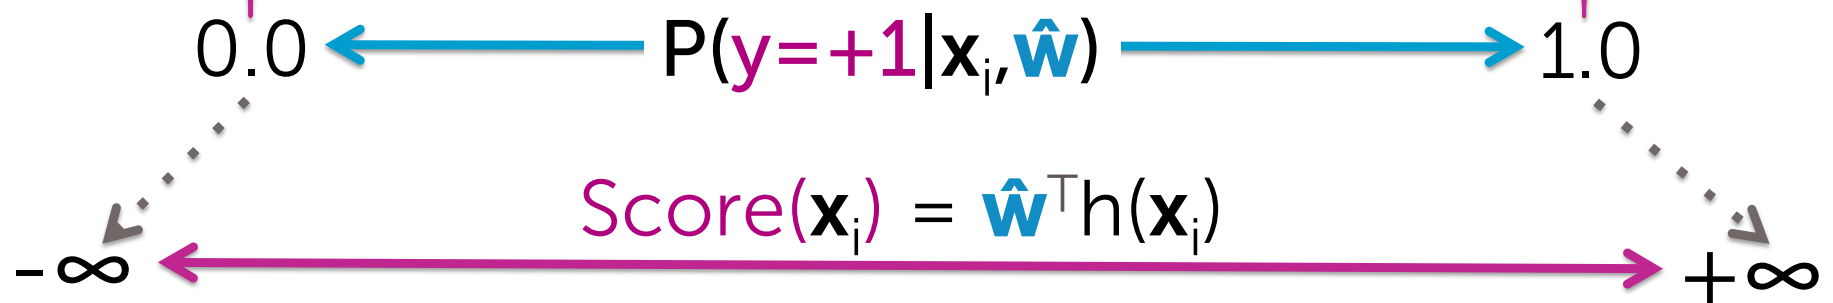


But,  $\text{sign}(z)$  only outputs -1 or +1,  
no probabilities in between

# Finding best coefficients

x[1] = #awesome	x[2] = #awful	y = sentiment
0	2	-1
3	3	-1
2	4	-1
0	3	-1
0	1	-1

x[1] = #awesome	x[2] = #awful	y = sentiment
2	1	+1
4	1	+1
1	1	+1
2	1	+1



# Quality metric: probability of data

$$\hat{P}(y=+1|\mathbf{x}, \hat{\mathbf{w}}) = \frac{1}{1 + e^{-\hat{\mathbf{w}}^T \mathbf{h}(\mathbf{x})}}$$

x[1] = #awesome	x[2] = #awful	y = sentiment
2	1	+1



If model good, should predict



Increase probability  $y=+1$  when



Choose  $\mathbf{w}$  to make

x[1] = #awesome	x[2] = #awful	y = sentiment
0	2	-1



If model good, should predict



Increase probability  $y=-1$  when



Choose  $\mathbf{w}$  to make

## Maximizing likelihood (probability of data)

Data point	$x[1]$	$x[2]$	$y$	Choose $w$ to maximize
$\mathbf{x}_1, y_1$	2	1	+1	
$\mathbf{x}_2, y_2$	0	2	-1	
$\mathbf{x}_3, y_3$	3	3	-1	
$\mathbf{x}_4, y_4$	4	1	+1	
$\mathbf{x}_5, y_5$	1	1	+1	
$\mathbf{x}_6, y_6$	2	4	-1	
$\mathbf{x}_7, y_7$	0	3	-1	
$\mathbf{x}_8, y_8$	0	1	-1	
$\mathbf{x}_9, y_9$	2	1	+1	

Must combine into single measure of quality



## Learn logistic regression model with maximum likelihood estimation (MLE)

- Choose coefficients  $\mathbf{w}$  that maximize likelihood:

$$\prod_{i=1}^N P(y_i \mid \mathbf{x}_i, \mathbf{w})$$

- No closed-form solution → use gradient ascent